

# Complete Code Guide: Customer Churn Prediction

---

## Introduction and Overview

## Code Explanation and Logic

---

Code Block 1:

# importing libraries and loading data set

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
file_path = '/content/WA_Fn-UseC_-Telco-Customer-Churn.csv'
df = pd.read_csv(file_path)
```

Explanation:

This block performs the following operations:

Code Block 2:

```
df.head()
```

Explanation:

This block performs the following operations:

Code Block 3:

```
df.info()  
df.describe()
```

Explanation:

This block performs the following operations:

Code Block 4:

```
statistics = df.describe(include='all')  
missing_values = df.isnull().sum()
```

statistics, missing\_values

Explanation:

This block performs the following operations:

Code Block 5:

```
data = {  
    'customerID': ['7590-VHVEG', '5575-GNVDE', '3668-QPYBK', '7795-CFOCW', '9237-  
HQUITU'],  
    'gender': ['Female', 'Male', 'Male', 'Male', 'Female'],  
    'SeniorCitizen': [0, 0, 0, 0, 0],  
    'Partner': ['Yes', 'No', 'No', 'No', 'No'],  
    'Dependents': ['No', 'No', 'No', 'No', 'No'],  
    'tenure': [1, 34, 2, 45, 2],  
    'PhoneService': ['No', 'Yes', 'Yes', 'No', 'Yes'],  
    'MultipleLines': ['No phone service', 'No', 'No', 'No phone service', 'No'],  
    'InternetService': ['DSL', 'DSL', 'DSL', 'DSL', 'Fiber optic'],  
    'OnlineSecurity': ['No', 'Yes', 'Yes', 'No', 'No'],  
    'OnlineBackup': ['Yes', 'No', 'Yes', 'No', 'Yes'],  
    'DeviceProtection': ['No', 'Yes', 'No', 'Yes', 'No'],  
    'TechSupport': ['No', 'No', 'No', 'Yes', 'No'],  
    'StreamingTV': ['No', 'No', 'No', 'Yes', 'Yes'],  
    'StreamingMovies': ['No', 'No', 'No', 'No', 'No'],
```

```

    'Contract': ['Month-to-month', 'One year', 'Month-to-month', 'One year', 'Month-to-month'],
    'PaperlessBilling': ['Yes', 'No', 'Yes', 'No', 'Yes'],
    'PaymentMethod': ['Electronic check', 'Mailed check', 'Mailed check', 'Bank transfer (automatic)', 'Electronic check'],
    'MonthlyCharges': [29.85, 56.95, 53.85, 42.30, 70.70],
    'TotalCharges': ['29.85', '1889.5', '108.15', '1840.75', '151.65'],
    'Churn': ['No', 'No', 'Yes', 'No', 'Yes']
}

```

Explanation:

This block performs the following operations:

Code Block 6:

```

df_sample = pd.DataFrame(data)
df_sample['TotalCharges'] = pd.to_numeric(df_sample['TotalCharges'], errors='coerce')

df_sample_info = df_sample.info()
df_sample_statistics = df_sample.describe(include='all')
df_sample_missing_values = df_sample.isnull().sum()
df_sample_info, df_sample_statistics, df_sample_missing_values

```

Explanation:

This block performs the following operations:

Code Block 7:

```

df.hist(bins=30, figsize=(20, 15))
plt.show()

plt.figure(figsize=(20, 15))
sns.boxplot(data=df)
plt.show()

```

Explanation:

This block performs the following operations:

Code Block 8:

```

df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')

```

```
df['TotalCharges'].fillna(df['TotalCharges'].median(), inplace=True)
```

Explanation:

This block performs the following operations:

Code Block 9:

```
Q1 = df['MonthlyCharges'].quantile(0.25)
Q3 = df['MonthlyCharges'].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
df = df[(df['MonthlyCharges'] >= lower_bound) & (df['MonthlyCharges'] <= upper_bound)]
```

Explanation:

This block performs the following operations:

Code Block 10:

```
df['tenure_years'] = df['tenure'] / 12
df = pd.get_dummies(df, drop_first=True)
```

Explanation:

This block performs the following operations:

Code Block 11:

```
scaler = StandardScaler()
numerical_features = ['tenure', 'MonthlyCharges', 'TotalCharges']
df[numerical_features] = scaler.fit_transform(df[numerical_features])
```

Explanation:

This block performs the following operations:

Code Block 12:

```
X = df.drop('Churn_Yes', axis=1)
y = df['Churn_Yes']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Explanation:

This block performs the following operations:

Code Block 13:

```
# random forest classifier
rf = RandomForestClassifier()
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
```

Explanation:

This block performs the following operations:

Code Block 14:

```
# Evaluate random forest classifier
```

```
metrics_rf = {
    "accuracy": accuracy_score(y_test, y_pred_rf),
    "precision": precision_score(y_test, y_pred_rf),
    "recall": recall_score(y_test, y_pred_rf),
    "f1_score": f1_score(y_test, y_pred_rf)
}
metrics_rf
```

Explanation:

This block performs the following operations:

Code Block 15:

```
# Logistic Regression
lr = LogisticRegression()
lr.fit(X_train, y_train)
y_pred_lr = lr.predict(X_test)
```

Explanation:

This block performs the following operations:

Code Block 16:

```
# Evaluate Logistic Regression model
metrics_lr = {
    "accuracy": accuracy_score(y_test, y_pred_lr),
    "precision": precision_score(y_test, y_pred_lr),
    "recall": recall_score(y_test, y_pred_lr),
    "f1_score": f1_score(y_test, y_pred_lr)
}

metrics_lr
```

### Explanation of the Metrics:

#### 1. Accuracy:

- **Linear Regression Accuracy:** 0.8226 (~82.3%)
- **Random Forest Accuracy:** 0.8048 (~80.5%)

Accuracy is the proportion of correct predictions out of the total number of predictions. In this case, Linear Regression has a slightly better accuracy (82.3%) compared to Random Forest (80.5%), meaning Linear Regression correctly predicted more instances overall.

#### 2. Precision:

- **Linear Regression Precision:** 0.6892 (~68.9%)
- **Random Forest Precision:** 0.6929 (~69.3%)

Precision measures how many of the predicted positive cases were actually positive. Both models have similar precision (~69%), indicating they are equally good at avoiding false positives (incorrectly predicting something as positive when it's actually negative).

#### 3. Recall:

- **Linear Regression Recall:** 0.6005 (~60.1%)
- **Random Forest Recall:** 0.4718 (~47.2%)

Recall measures how many of the actual positive cases were correctly predicted. Linear Regression has a significantly better recall (~60%) compared to Random Forest (~47.2%),

meaning Linear Regression is better at identifying true positives, while Random Forest misses more actual positives.

#### 4. F1 Score:

- **Linear Regression F1 Score:** 0.6418 (~64.2%)
- **Random Forest F1 Score:** 0.5614 (~56.1%)

The F1 score is the harmonic mean of precision and recall, providing a balance between the two. Linear Regression has a better F1 score (64.2%) compared to Random Forest (56.1%), indicating that Linear Regression strikes a better balance between precision and recall.

#### What This Means:

- **Linear Regression:**

- This model has a higher **accuracy** (82.3%), meaning it makes more correct predictions overall.
- The **precision** (68.9%) is slightly lower than Random Forest, but still comparable, meaning it is fairly reliable when it predicts something as positive.
- The **recall** (60.1%) is better, indicating that it correctly identifies a higher number of actual positives compared to Random Forest.
- Its **F1 score** (64.2%) reflects a better balance between precision and recall.

- **Random Forest:**

- This model has slightly lower **accuracy** (80.5%) than Linear Regression.
- It has a marginally higher **precision** (69.3%), meaning it is slightly better at avoiding false positives.
- However, the **recall** (47.2%) is lower, indicating that Random Forest misses more true positives than Linear Regression.
- The **F1 score** (56.1%) is lower, showing that it doesn't balance precision and recall as well as Linear Regression.

#### Summary:

- **Linear Regression** seems to perform better overall, particularly in terms of recall and F1 score, indicating it is better at identifying true positives and balancing precision and recall.

- **Random Forest** has a slightly better precision, which means it's more careful in avoiding false positives, but it sacrifices recall, missing more true positives.